

# Creating a model based on MNIST Dataset of Handwritten Digits

- Lets Load our Dataset

```
In [2]: from keras.datasets import mnist  
  
#Load MNIST Dataset  
  
(x_train, y_train), (x_test,y_test) = mnist.load_data()  
  
print(y_train.shape)
```

Using TensorFlow backend.

(60000,)

## Step 2A: Examine the size and image dimensions

- Check the number of samples, dimensions and whether images are color or gray
- We see that our training data consist of 60000 samples of training data, and 10000 samples of test data
- Our labels are properly sized as well
- Our Image dimension are 28x28 with no color channels.

In [3]: *# printing number of training samples and test samples*

```
print("Initial Shape of x_train", str(x_train.shape))

print("Number of samples in our training dataset:" + str(len(x_train)))
print("Number of labels in our training dataset:" + str(len(y_train)))
print("Number of samples in our test dataset:" + str(len(x_test)))
print("Number of labels in our test dataset:" + str(len(y_test)))

print()

print("Dimension of x_train:" + str(x_train[0].shape))
print("Labels in x_train:" + str(y_train.shape))

print()

print("Dimension of x_test:" + str(x_test[0].shape))
print("Labels in x_test:" + str(y_test.shape))
```

Initial Shape of x\_train (60000, 28, 28)  
Number of samples in our training dataset:60000  
Number of labels in our training dataset:60000  
Number of samples in our test dataset:10000  
Number of labels in our test dataset:10000

Dimension of x\_train:(28, 28)  
Labels in x\_train:(60000,)

Dimension of x\_test:(28, 28)  
Labels in x\_test:(10000,)

## Step 2B - Let's take a look at some of images in the dataset

- Using OpenCV
- Using Matplotlib

```
In [4]: #using OpenCV
#import OpenCV and numpy
import cv2
import numpy as np

#Using OpenCV to display 6 random images from our dataset

for i in range(0,6):
    random_num = np.random.randint(0, len(x_train))
    img = x_train[random_num]
    window_name = 'Random Sample #' + str(i)
    cv2.imshow(window_name, img)
    cv2.waitKey(0)

cv2.destroyAllWindows()
```

## Lets do the same thing using matplotlib to plot 6 images

```
In [4]: #import matplotlib
import matplotlib.pyplot as plt

#plt 6 images, subplots arguments are nrows, ncols, index
#we set the color map to grey since our dataset is grayscale

plt.subplot(331)
random_num = np.random.randint(0, len(x_train))
plt.imshow(x_train[random_num], cmap=plt.get_cmap('gray'))

plt.subplot(332)
random_num = np.random.randint(0, len(x_train))
plt.imshow(x_train[random_num], cmap=plt.get_cmap('gray'))


plt.subplot(333)
random_num = np.random.randint(0, len(x_train))
plt.imshow(x_train[random_num], cmap=plt.get_cmap('gray'))

plt.subplot(334)
random_num = np.random.randint(0, len(x_train))
plt.imshow(x_train[random_num], cmap=plt.get_cmap('gray'))

plt.subplot(335)
random_num = np.random.randint(0, len(x_train))
plt.imshow(x_train[random_num], cmap=plt.get_cmap('gray'))

plt.subplot(336)
random_num = np.random.randint(0, len(x_train))
plt.imshow(x_train[random_num], cmap=plt.get_cmap('gray'))

#Display out plots
plt.show()
```



## Step 3A - Prepare our dataset for training

```

In [5]: #Lets Store the number of rows and columns
img_rows = x_train[0].shape[0]
img_cols = x_train[1].shape[0]

#Creating our data in right 'shape' needed for keras
#We need to add a 4th dimension to our data which is depth
#our original image shape is (60000,28,28) we will add depth to it (60000,28,28,1)
x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols,1)
x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols,1)

#store the shape of a single image
#here the input shape is in integer
#we created image shape for our first layer of convolution (1 because of grayscale)
input_shape = (img_rows, img_cols, 1)

#Change our image type to float32 data type
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')

#Normalize our data from a range of (0-255) to (0-1)
x_train /= 255
x_test /= 255

#Printing our modified data in accordance with keras
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

```

```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

## Step 3B - Hot Encode our labels(Y)

```

In [6]: from keras.utils import np_utils

#Now we hot one encoded outputs
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)

#Count the number of columns in our hot encoded matrix
print ('Number of Classes:' + str(y_test.shape[1]))

num_classes = y_test.shape[1]
num_pixels = x_train.shape[1] * x_train.shape[2]

```

```

Number of Classes:10

```

## Step 4 - Create Our Model

- We're constructing a simple CNN that uses 32 filters of size 3\*3
- We've added a second convolution layer of 64 filters of the same size 3\*3

- We downsampled our data to  $2 \times 2$ , here we apply a dropout where  $p$  is set to 0.25
- After that we flattened our Max Pool output that is connected to a Dense/FC layer that has an output size of 128
- In the final step 128 output is connected to another FC/Dense layer that outputs to the 10 categorical units
- Note dropout shouldn't be greater than 0.5 preferably

```
In [7]: import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
from keras.optimizers import SGD

#Create model
model = Sequential()

model.add(Conv2D(32, kernel_size=(3,3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64,(3,3),activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128,activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss = 'categorical_crossentropy',
              optimizer = SGD(0.01),
              metrics = ['accuracy'])

print(model.summary())
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
flatten_1 (Flatten)	(None, 9216)	0
dense_1 (Dense)	(None, 128)	1179776
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290
=====		
Total params: 1,199,882		
Trainable params: 1,199,882		
Non-trainable params: 0		
None		

## Step 5 - Train our Model

- we enter our formatted data as inputs and set the batch size, number of epochs
- we store our models training result for plotting in future
- Then we used keras model evaluation function to output the model's final performance. Here we are examining Test Loss and Test Accuracy

```
In [9]: batch_size = 32
epochs = 10

history = model.fit(x_train,
                    y_train,
                    batch_size = batch_size,
                    epochs = epochs,
                    verbose = 1,                                #how much data
                    validation_data = (x_test, y_test))

score = model.evaluate(x_test, y_test, verbose = 0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 115s 2ms/step - loss: 0.5812 - accuracy: 0.8182 - val\_loss: 0.1976 - val\_accuracy: 0.9393- accuracy

Epoch 2/10

60000/60000 [=====] - 117s 2ms/step - loss: 0.2949 - accuracy: 0.9112 - val\_loss: 0.1417 - val\_accuracy: 0.9582

Epoch 3/10

60000/60000 [=====] - 119s 2ms/step - loss: 0.2260 - accuracy: 0.9322 - val\_loss: 0.1074 - val\_accuracy: 0.9664

Epoch 4/10

60000/60000 [=====] - 122s 2ms/step - loss: 0.1761 - accuracy: 0.9474 - val\_loss: 0.0836 - val\_accuracy: 0.9745

Epoch 5/10

60000/60000 [=====] - 119s 2ms/step - loss: 0.1461 - accuracy: 0.9560 - val\_loss: 0.0702 - val\_accuracy: 0.9776 - loss: 0.1462 - accuracy

Epoch 6/10

60000/60000 [=====] - 119s 2ms/step - loss: 0.1248 - accuracy: 0.9629 - val\_loss: 0.0592 - val\_accuracy: 0.9811

Epoch 7/10

60000/60000 [=====] - 121s 2ms/step - loss: 0.1101 - accuracy: 0.9677 - val\_loss: 0.0537 - val\_accuracy: 0.9825

Epoch 8/10

60000/60000 [=====] - 119s 2ms/step - loss: 0.0999 - accuracy: 0.9695 - val\_loss: 0.0501 - val\_accuracy: 0.9838

Epoch 9/10

60000/60000 [=====] - 119s 2ms/step - loss: 0.0915 - accuracy: 0.9726 - val\_loss: 0.0452 - val\_accuracy: 0.9840

Epoch 10/10

60000/60000 [=====] - 121s 2ms/step - loss: 0.0839 - accuracy: 0.9745 - val\_loss: 0.0438 - val\_accuracy: 0.9863

Test loss: 0.043802500181901266

Test accuracy: 0.986299991607666

## Step 6 - Plotting our Loss and Accuracy Charts


In [10]: *# Plotting our Loss Charts*

```
import matplotlib.pyplot as plt

history_dict = history.history

loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']
epochs = range(1, len(loss_values) + 1)

line1 = plt.plot(epochs, val_loss_values, label='Validation/Test Loss')
line2 = plt.plot(epochs, loss_values, label='Training Loss')
plt.setp(line1, linewidth=2.0, marker = '+', markersize=10.0)
plt.setp(line2, linewidth=2.0, marker = 'x', markersize=10.0)
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid(True)
plt.legend()
plt.show()
```






In [12]: *#plotting accuracy charts*

```
import matplotlib.pyplot as plt

history_dict = history.history

accuracy_values = history_dict['accuracy']
val_accuracy_values = history_dict['val_accuracy']
epochs = range(1, len(accuracy_values) + 1)

line1 = plt.plot(epochs, val_accuracy_values, label='Validation/TestAccuracy')
line2 = plt.plot(epochs, accuracy_values, label='Training Accuracy')
plt.setp(line1, linewidth=2.0, marker = '+', markersize=10.0)
plt.setp(line2, linewidth=2.0, marker = '4', markersize=10.0)
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.grid(True)
plt.legend()
plt.show()
```



## Step 7A - Saving our Model

```
In [13]: model.save(r"C:\Users\yrsin\Desktop\git_Hub\tensorflow\sharedwindows10\testfolder")
print("model Saved")
```

model Saved

## Step 7B - Loading our Model

```
In [14]: from keras.models import load_model

classifier = load_model(r'C:\Users\yrsin\Desktop\git_Hub\tensorflow\sharedwindows10\testfolder')
```

## Step 8 - Lets input some data to test our

# classifier

```
In [17]: import cv2
import numpy as np

def draw_test(name, pred, input_im):
    BLACK = [0,0,0]
    expanded_image = cv2.copyMakeBorder(input_im, 0,0,0, imageL.shape[0], cv2.BORDER_CONSTANT, BLACK)
    expanded_image = cv2.cvtColor(expanded_image, cv2.COLOR_GRAY2BGR)
    cv2.putText(expanded_image, str(pred), (152,70), cv2.FONT_HERSHEY_COMPLEX_SMALL, 1, (0,255,0), 2)
    cv2.imshow(name, expanded_image)

for i in range(0,10):
    rand = np.random.randint(0,len(x_test))
    input_im = x_test[rand]

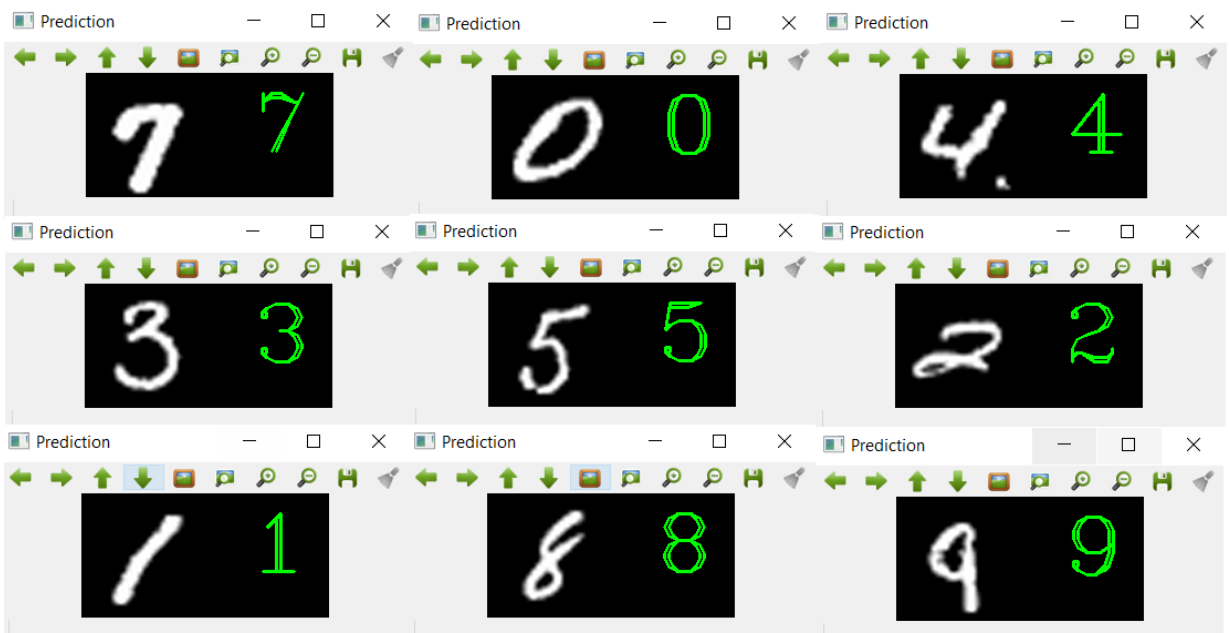
    imageL = cv2.resize(input_im, None, fx=4, fy=4, interpolation = cv2.INTER_CUBIC)
    input_im = input_im.reshape(1,28,28,1)

    #Get Prediction
    res = str(classifier.predict_classes(input_im,1,verbose = 0)[0])

    draw_test("Prediction",res,imageL)
    cv2.waitKey(0)

cv2.destroyAllWindows()
```

## Predictions of Classifier



## Confusion Matrix and classification Report

```
In [19]: from sklearn.metrics import classification_report, confusion_matrix
import numpy as np
```

```
y_pred = model.predict_classes(x_test)
```

```
print(classification_report(np.argmax(y_test,axis=1), y_pred))
```

```
print(confusion_matrix(np.argmax(y_test,axis=1), y_pred))
```

	precision	recall	f1-score	support
0	0.97	1.00	0.98	980
1	0.99	0.99	0.99	1135
2	0.98	0.99	0.99	1032
3	0.99	0.98	0.99	1010
4	0.99	0.99	0.99	982
5	0.98	0.99	0.99	892
6	0.99	0.98	0.99	958
7	0.99	0.98	0.98	1028
8	0.99	0.98	0.98	974
9	0.99	0.98	0.98	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

```
[[ 976    1    0    0    0    0    1    1    1    0]
 [    0 1127    3    1    0    1    2    0    1    0]
 [    3    2 1022    1    1    0    0    3    0    0]
 [    0    0    3  994    0    8    0    3    2    0]
 [    0    1    1    0  968    0    4    0    2    6]
 [    2    0    0    4    0  882    2    1    1    0]
 [    8    2    0    0    1    4  942    0    1    0]
 [    2    1   10    1    0    0    0 1011    1    2]
 [    7    1    2    1    2    1    1    3  952    4]
 [    5    4    0    2    3    1    0    4    1  989]]
```

## Displaying our Misclassified Data

```
In [20]: import cv2
import numpy as np
from keras.datasets import mnist

# Loads the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Use numpy to create an array that stores a value of 1 when a misclassification
result = np.absolute(y_test - y_pred)
result_indices = np.nonzero(result > 0)

# Display the indices of misclassifications
print("Indices of misclassified data are: \n\n" + str(result_indices))
```

Indices of misclassified data are:

```
(array([ 247,  259,  321,  445,  449,  582,  720,  726,  740,  844,  883,
        938,  947,  965, 1014, 1039, 1112, 1182, 1226, 1232, 1242, 1247,
       1290, 1299, 1319, 1326, 1393, 1522, 1527, 1530, 1549, 1609, 1678,
       1681, 1709, 1717, 1737, 1878, 1901, 2018, 2035, 2043, 2093, 2098,
       2109, 2118, 2130, 2135, 2182, 2272, 2293, 2387, 2406, 2447, 2454,
       2462, 2488, 2607, 2648, 2654, 2896, 2921, 2927, 2939, 2952, 2953,
       3030, 3060, 3289, 3422, 3503, 3520, 3558, 3559, 3597, 3604, 3726,
       3751, 3767, 3780, 3808, 3853, 3906, 3941, 4075, 4078, 4163, 4176,
       4205, 4224, 4238, 4248, 4256, 4265, 4289, 4497, 4500, 4536, 4601,
       4639, 4740, 4761, 4807, 4814, 4823, 4956, 5331, 5642, 5734, 5887,
       5937, 5955, 5973, 6011, 6505, 6555, 6572, 6576, 6597, 6625, 6651,
       6783, 7434, 8325, 8408, 8527, 9009, 9015, 9019, 9024, 9280, 9634,
       9692, 9729, 9770, 9888, 9982], dtype=int64),)
```

## Displaying the Misclassification

```

In [21]: import cv2
#from keras.models import load_model

#classifier = load_model('/home/deeplearningcv/DeepLearningCV/Trained Models/mnist

def draw_test(name, pred, input_im, true_label):
    BLACK = [0,0,0]
    expanded_image = cv2.copyMakeBorder(input_im, 0, 0, 0, imageL.shape[0]*2 ,cv2
    expanded_image = cv2.cvtColor(expanded_image, cv2.COLOR_GRAY2BGR)
    cv2.putText(expanded_image, str(pred), (152, 70) , cv2.FONT_HERSHEY_COMPLEX_S
    cv2.putText(expanded_image, str(true_label), (250, 70) , cv2.FONT_HERSHEY_COM
    cv2.imshow(name, expanded_image)

for i in range(0,10):

    input_im = x_test[result_indices[0][i]]
    #print(y_test[result_indices[0][i]])
    imageL = cv2.resize(input_im, None, fx=4, fy=4, interpolation = cv2.INTER_CUBI
    input_im = input_im.reshape(1,28,28,1)

    ## Get Prediction
    res = str(model.predict_classes(input_im, 1, verbose = 0)[0])
    draw_test("Prediction", res, imageL, y_test[result_indices[0][i]])
    cv2.waitKey(0)

cv2.destroyAllWindows()

```

## Output of Misclassified data

